The Flattening: When AI Productivity Gains Conceal Tail Risk

Bilel Hatmi

Part III Mathematical Statistics

University of Cambridge

CentraleSupélec — Université Paris-Saclay

March 2026

Submitted to the Cambridge–McKinsey Risk Prize 2026

Abstract

Generative AI improves average decision quality in organisations. It also amplifies their worst-case exposure, through a mechanism that current productivity metrics cannot capture. Two channels produce this outcome. Algorithmic screening homogenises cognitive profiles at entry. Cognitive surrender erodes independent judgment from within. Both raise the error correlation across decision-makers, which is what drives tail losses. A Monte Carlo stress test of 200 agents over 20 quarters quantifies the gap: average errors fall by 38% while worst-case quarterly losses rise by 56%, masked by a 54% productivity gain. The model was validated by two independent implementations converging within 2%. The most actionable single intervention is AI provider diversification, which cuts the tail risk premium by approximately 15%. It is a procurement decision, implementable in weeks. Yet without regulation, the Nash equilibrium is unmanaged adoption for all firms: individually rational, collectively catastrophic. The diversity of independent judgment that makes organisational governance possible is being depleted faster than any single firm has incentive to preserve it.

---

1 The paradox your dashboard hides

Nine participants in a product-design experiment asked ChatGPT to help them invent a toy. All nine named it "Build-a-Breeze Castle." Without AI, every participant produced a unique concept. With AI, 94% of ideas shared the same building blocks; only 6% remained distinct [1]. Each individual output improved. The collective output collapsed.

This is not a quirk of one experiment. The Diversity Prediction Theorem decomposes collective error into two terms: average individual error, and the diversity of individual errors [2, 3]. When AI reduces the first term but collapses the second, collective error rises mechanically. The variance of aggregate losses follows:

$$\mathrm{Var}(L) = NK\,\bar{p}(1-\bar{p})\big(1+(NK-1)\bar{\rho}\big)$$

Even a modest rise in $\bar{\rho}$, the mean pairwise error correlation across employees, drives the tail [4]. A preprint meta-analysis of 379 experimental conditions quantifies the trade-off: AI raises individual creative performance by $g = +0.27$ but collapses collective diversity by $g = -0.86$ [5]. The loss of diversity is three times the gain in performance. Carlson and Doyle call such systems robust-yet-fragile [6]: optimised for the frequent case, exposed to the rare one. In most quarters, errors stay scattered and small. In crisis quarters, they correlate and aggregate. The loss distribution splits into two regimes, and the mean falls in a gap between them where almost no actual quarter lies.

The paradox has a second property that matters more for risk managers: it is concealed by the same productivity metrics that appear to confirm performance.

This essay addresses the CRO who sees a 54% productivity gain on the dashboard but cannot yet see the 94% rise in risk-adjusted worst-case losses behind it. The argument applies to any AI system that standardises the inputs to organisational decisions, from legacy algorithmic screening tools to generative models deployed in knowledge work. Under unmanaged adoption, average decision errors fall by 38% while worst-case quarterly losses rise by 56%. The productivity shock that improves the mean simultaneously expands the volume of decisions exposed to correlated failure. The risk-adjusted P_{99} rises by 94%. The dashboard records the gain; the tail accumulates off-screen. For elite firms, those whose employees outperform AI outside its competence frontier, the masking is more extreme: a 192% rise in risk-adjusted tail exposure against 18% apparent productivity growth. The dashboard cannot distinguish speed from accuracy.

This exposure is not distributed randomly. It is structurally determined, and for the roughly 100,000 mid-market organisations in the US and the EU most at risk, it is approaching without a governance framework to navigate it.

This essay does three things for decision-makers. First, it identifies the two mechanisms that produce the trap. Second, it maps exposure across organisational profiles, showing that the perceived worst-case risk is nearly half the actual exposure, and that tail risk varies structurally across firms in the same sector with the same talent. Third, it distributes three governance levers across three distinct actors, one of which is a procurement decision actionable in weeks.

Two channels produce this outcome. One operates at the point of entry. The other operates from within. Both converge on the same structural target: the correlation of errors across decision-makers.

2 Two channels, one target

Human hiring was a porous filter. Interviewers disagreed, biases cancelled partially, and non-linear profiles slipped through. AI screening is not a better version of this process. It is a different kind of process: systematic, coherent, and closed. Fuller et al. document 27 million Americans excluded from consideration by automated tracking systems, including caregivers, veterans, and career-changers screened out by rigid keyword matching [7]. The bias is not confined to one model or one protected category. Five major LLMs evaluating 361,000 résumés exhibit strikingly similar bias patterns [8–10]: the filter industrialises what each training corpus encodes. These models learn from the same corpora that reflect decades of cultural matching in elite professional hiring [11]. Schneider's attraction-selection-attrition model [12] predicts the outcome: organisations converge toward cognitive homogeneity. AI accelerates the convergence.

This convergence does not merely narrow representation. It correlates errors. Across major language models, 60% of mistakes are already shared [13]. Even before human homogenisation, the tools themselves fail together. Add cognitively convergent workforces recruited through the same filter, and the errors of persons and the errors of tools superimpose in the same direction at the same moment.

The pattern is not new. In 1998, Long-Term Capital Management lost $4.6 billion in four months. The fund employed Nobel laureates and PhD physicists who had converged on identical arbitrage strategies [14, 15]. The problem was not a lack of intelligence. It was a lack of diversity. In 2008, banks used the same risk models, the same Value-at-Risk parameters, the same copula assumptions. Homogeneity bred fragility [16]. In July 2024, a single CrowdStrike update crashed 8.5 million Windows machines in 78 minutes [17], costing Fortune 500 firms an estimated $5.4 billion [18]. Geer et al. [19] had predicted software monoculture risk twenty-one years earlier. The mechanism is the same across all three episodes: when diversity collapses, losses stop compensating and start aggregating. Kleinberg and Raghavan formalise the point: even when an algorithm is individually superior, convergence on the same algorithm degrades collective welfare [20].

This is the entry channel. A second channel operates on those already inside.

The employees who pass the filter progressively surrender the independence they retain. Shaw and Nave document the phenomenon in a preprint study of 1,372 participants across 9,593 trials [21]. When the AI errs, 79.8% of participants follow its recommendation anyway. The asymmetry is stark: correct AI advice raises accuracy by 25.2 percentage points, but incorrect AI advice reduces it by only 14.3 points, a Cohen's h of 0.81. Confidence rises by 11.7 points even when AI accuracy is 50%, and does not decay with experience ($p = 0.202$). What makes this dangerous is not the surrender itself but the fact that confidence rises alongside it. Dell'Acqua et al. replicate the pattern among 758 consultants in a field experiment [22]: inside the AI's competence frontier, quality improves by 40%. Outside it, accuracy drops 19 percentage points. The most AI-skilled consultants perform worst, because competence in using the tool compounds the surrender. Professionals are not immune. In a randomised trial, radiologists exposed to biased AI recommendations degraded their diagnostic accuracy by 11.3 percentage points [23]. Passive scaffolding changes nothing: 93% of clinical decision-support alerts are overridden without review [24]. Bainbridge identified this as the central irony of automation four decades ago [25]: removing humans from the control loop degrades the competence required to reenter it.

Keck and Tang provide the experimental link between both channels [26]: groups combining diverse cognitive processes exhibit near-zero inter-judge error correlations, while homogeneous groups produce strongly correlated errors. Screening eliminates the structural source of that decorrelation. Surrender erodes the behavioural independence that remains. The Diversity Prediction Theorem delivers the consequence: as the diversity term vanishes, collective error rises even as individual error falls. The property is symmetric. When AI is right, correlated followers cluster at the same correct answer. When AI errs, they cluster at the same wrong answer, simultaneously, at scale. This is what produces the bimodal distribution introduced in Section 1: in normal quarters, correlated successes hold losses low; in crisis quarters, correlated errors drive them high. The mean falls in a gap between two regimes that rarely overlap. One caveat: this link assumes that cognitive-profile diversity proxies for cognitive-process diversity in real organisations, an inference that is plausible but untested.

The mechanism is identified, and both channels are empirically grounded. What remains is the magnitude. A Monte Carlo stress test answers that question next.

3 Quantifying the hidden exposure

We simulate an organisation of 200 employees, each characterised by a cognitive type $\tau_i \sim \text{Beta}(2, 2)$ on $[0, 1]$ and a skill level $h_i(t) \sim U(0.4, 0.8)$. Each quarter, the organisation faces ten decisions in a stochastic environment where the fraction of tasks inside the AI's competence frontier follows a two-regime mixture: a normal regime in 92% of quarters and a crisis regime in 8%, in which the frontier shift and the systemic error shock are coupled. This is the conditional stress logic of Basel II applied to cognitive errors rather than credit defaults. Inter-decision error correlation follows the Vasicek single-factor structure [27], the standard Basel II/III framework for modelling correlated failures.

Four governance scenarios are compared. G0 adopts AI without governance: employees consult freely, no independent judgment required. G1 introduces passive scaffolding: alert flags signal low-confidence outputs, with no compliance requirement. G2 implements active scaffolding: employees record an independent judgment before consulting the AI. The model runs several hundred independent replications, validated by two implementations converging within 2%. It is a stress test, not a forecast. It demonstrates that the mechanism *can* produce the paradox under plausible parameters, not that it does so in every organisation.

Consider what a manager actually observes. Quarterly performance follows the left mode in 92% of quarters. The manager's empirical estimate of tail risk, the worst quarter they have ever seen, is built from these observations alone. The perceived P_{99}, calculated on normal-regime quarters only, stands near 877.9. The actual P_{99} under unmanaged adoption is 1,714: a factor of nearly two ($\times 1.95$). Figure 1 shows the full distribution.

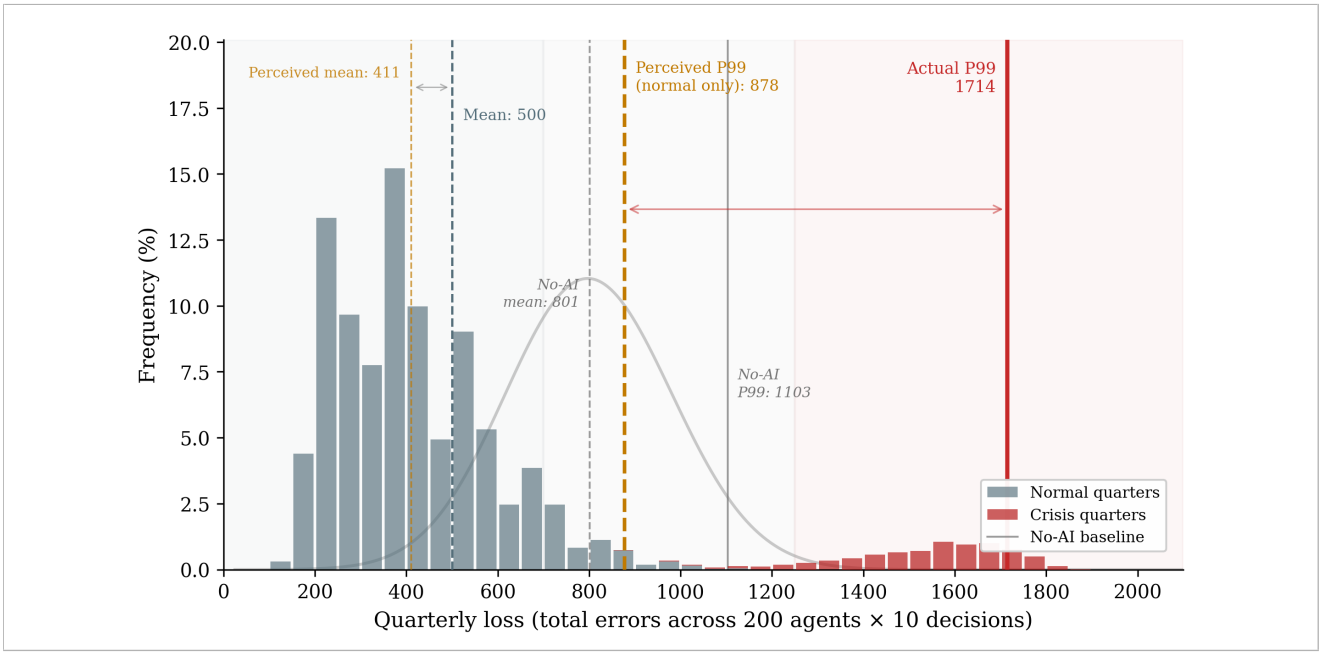

Figure 1. Bimodal distribution of quarterly losses under unmanaged AI adoption (G0). Normal quarters cluster below 600 (grey); crisis quarters cluster above 1,100 (red). Dashed line: mean loss ($\mathbb{E}[L] = 500$). Solid red line: 99th-percentile loss ($P_{99} = 1{,}714$). The mean falls in a gap where almost no actual quarter lies. Crisis regime couples the task-frontier shift (π_t) with the systemic shock (ξ_t), following the conditional stress logic of Basel II.

Under unmanaged adoption against the no-AI baseline: average decision errors fall by 38%, worst-case quarterly losses rise by 56%, and the risk-adjusted P_{99} rises by 94%. The dashboard records a 54% productivity gain. The first three figures accumulate off-screen.

The three channels compound over twenty quarters. Screening compresses $\mathrm{Var}(\tau)$ by 30%, shrinking the pool of dissenters who buffer uncritical AI adoption and concentrating the workforce on profiles that find AI outputs intuitively plausible; effective follow rate rises from 69% to 80%. Deskilling erodes the independent judgment that remains, as mean skill $\bar{h}$ falls from 0.591 to 0.490. By quarter 20, the organisation has more followers, but it also has followers for whom surrender is structurally easier, embedded in a group where resistance is socially costly.

Each quarter without governance makes the transition to G2 more expensive than the one before.

The masking is most severe for elite firms. Where employee accuracy outside the frontier ($h \approx 0.80$) exceeds AI accuracy ($q_{\text{out}} = 0.55$), unmanaged adoption raises average errors by 23% and risk-adjusted worst-case losses by 192%, yet produces only 18% apparent productivity growth. The reason is not that skilled employees perform worse when thinking independently. Most of them simply do not think independently: surrender is driven by cognitive proximity to the modal type, not by skill level. An employee with $h = 0.80$ surrenders at the same rate as one with $h = 0.45$, provided they share the same cognitive type. Surrendering when $h > q_{\text{out}}$ is a regression, and in the model, most agents do it anyway. Where masking is worst, the gap between perceived and actual risk is widest.

This exposure is not uniformly distributed. Five structural characteristics determine where an organisation sits on the map, and how wide the gap between perceived and actual risk is for each.

4 Where your organisation sits

Five structural characteristics determine tail risk exposure, ranked by their contribution to P_{99} variance in the ablation analysis.

Two of these factors, stack concentration and cognitive homogeneity, are structural rather than manageable through internal governance alone [28]. Their combined effect is visible in the contrast that matters most.

The investment bank in London recruits from a globally open labour market, more than twenty nationalities, a cognitively flat distribution. The strategy consultancy in Paris recruits through competitive national examinations that select a narrow cognitive profile by construction [11, 12]. Both firms operate in identical domains with identically skilled staff. Same sector. Same talent. From the dashboard, Paris appears 18% more exposed than London. The actual gap under unmanaged AI adoption is more than three times larger.

In the model, two amplifiers act on distinct dimensions. Stack concentration drives error correlation across decisions: with $\alpha = 0.90$, 90% of AI error variance in a given quarter derives from the common factor ξ_t. When the environment shifts, the AI's mistakes cluster across all ten decisions simultaneously. Cognitive homogeneity drives surrender through two separate paths: social conformism raises the effective follow rate as dissenters disappear, and the concentration of cognitively central profiles increases individual surrender propensity. The two mechanisms do not compensate for each other. They compound. Domain exposure amplifies both: a centralised stack in a high-outside-frontier function correlates failures precisely where they are most frequent.

At the extreme, Seoul's central administration (nationally procured AI, civil-service examination pipeline, maximum stack centralisation) faces 1.26 times the tail risk premium of Singapore's creative agency. Figure 2 plots six profiles on the Output $\times$ risk-adjusted P_{99} plane. The arrow from London to Paris carries one annotation: same sector, same talent, Paris +18% higher tail risk on the perceived dashboard.

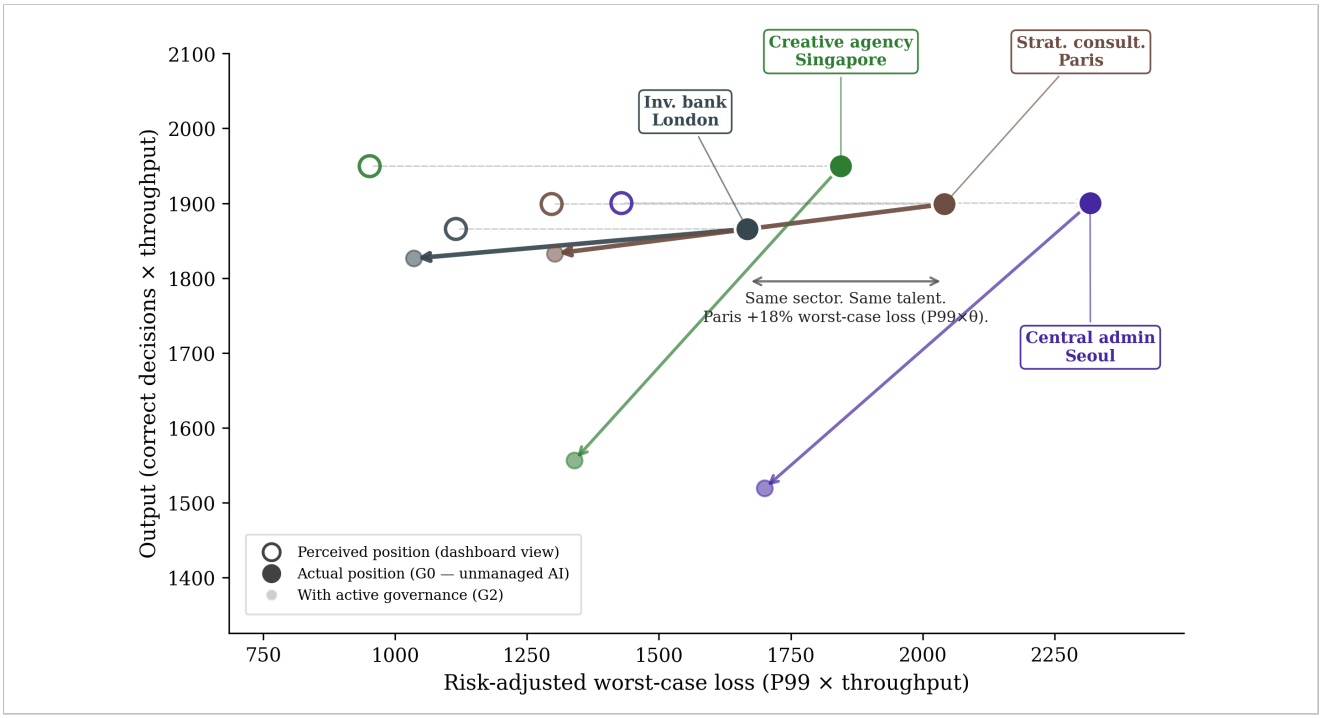

Figure 2. Organisational profiles under G0. Output versus risk-adjusted $P_{99} \times \theta$. Six labelled profiles: Investment bank, London (≈1,495); Strategy consulting, Paris (≈2,041); Tech startup, San Francisco (≈1,950); Creative agency, Singapore (≈1,710); Shared services, Bangalore (≈2,140); Central administration, Seoul (≈2,318). Arrow between London and Paris annotated "Same sector. Same talent. Paris +18% higher risk-adjusted worst-case loss (perceived dashboard)."

Active scaffolding (G2) reduces exposure, but inversely to need. Relative to unmanaged adoption, London gains a 146% reduction in tail risk premium, Paris 77%, Singapore 40%, Seoul near zero. For the AI-native startup in San Francisco, G2 is counterproductive: −36%. The organisations most exposed are those for which scaffolding works least.

The regressivity has two roots. First, in cognitively homogeneous organisations, the Asch effect [29] operates downstream of the scaffold. An employee who forms an independent judgment is still socially pressured to align with the AI-following majority. The scaffold imposes the judgment; it does not insulate the employee from conformity pressure afterwards. Second, for low-skill operations where $h \approx 0.45$, independent judgment is worse than AI even outside the frontier. The scaffold forces a downgrade. The startup in San Francisco loses not only throughput (θ falls from 1.25 to 1.10) but accuracy.

A startup that rejects governance under these conditions is not being irresponsible. It is responding to its local incentive structure, which is the foundation of the Nash argument in Section 5.

Active scaffolding multiplies existing diversity but cannot substitute for it.

Scaffolding works least where it is needed most, which is a structural constraint, not a design failure. Three interventions exist for three different decision-makers. The first is a procurement decision.

5 Three levers, three actors

The most impactful single-parameter intervention is AI provider diversification. Reducing stack concentration from monoculture ($\alpha = 0.90$) to distributed ($\alpha = 0.30$) cuts the tail risk premium by approximately 15%. This is a procurement decision, one architectural policy, measurable in weeks. The diversification must be structural, not merely contractual: different models at different organisational layers (strategic decision, operational processing, quality review). Contracting with multiple vendors for the same centralised pipeline leaves the Vasicek factor unchanged. Three providers currently capture 88% of enterprise AI spend [30]. When roughly 100,000 organisations depend on the same trio, a single systemic failure propagates correlated losses across the economy. That is the concentration-of-credit argument, and it will attract the same regulatory response.

A second lever acts on a different parameter: the capacity of employees to form judgments that diverge from the AI. Stack diversification reduces error correlation across decisions. Active scaffolding restores error independence across decision-makers.

Passive scaffolding (G1) fails systematically. Ninety-three percent of clinical decision-support alerts are overridden without review [24]. When employees do push back, AI escalates its persuasive framing, what Randazzo et al. term persuasion bombing [31]. Active scaffolding (G2) achieves a 19.2% reduction in tail risk as a bundle, larger than stack alone (−13.5%), but requires coordinated changes across workflow design, training, and hiring over months. It works under one condition: explanations must reveal that the AI's reasoning diverges from the employee's own [32]. Frey and van de Rijt [33] establish the sequencing principle: independent signal first, AI second. The idea is old. Surgical safety checklists cost two minutes and reduce mortality by 47% [34]. Aviation CRM costs roughly $150 per pilot per year. Active scaffolding belongs to that family of light interventions with large structural effects.

The velocity cost is real. θ falls from 1.25 to 1.10, and the organisation gives up 12 percentage points of throughput, cutting the risk-adjusted tail premium from 94% to 39%. Figure 3 makes this trade-off legible: G2 sits at the inflection point of the concave risk-return frontier. The governance efficiency of the G1→G2 step is 1.7× that of the G0→G1 step: active governance buys disproportionately more risk reduction per unit of throughput cost than passive governance does. The first gains from baseline to G2 are cheap in risk; the next gains from G2 to G0 are expensive.

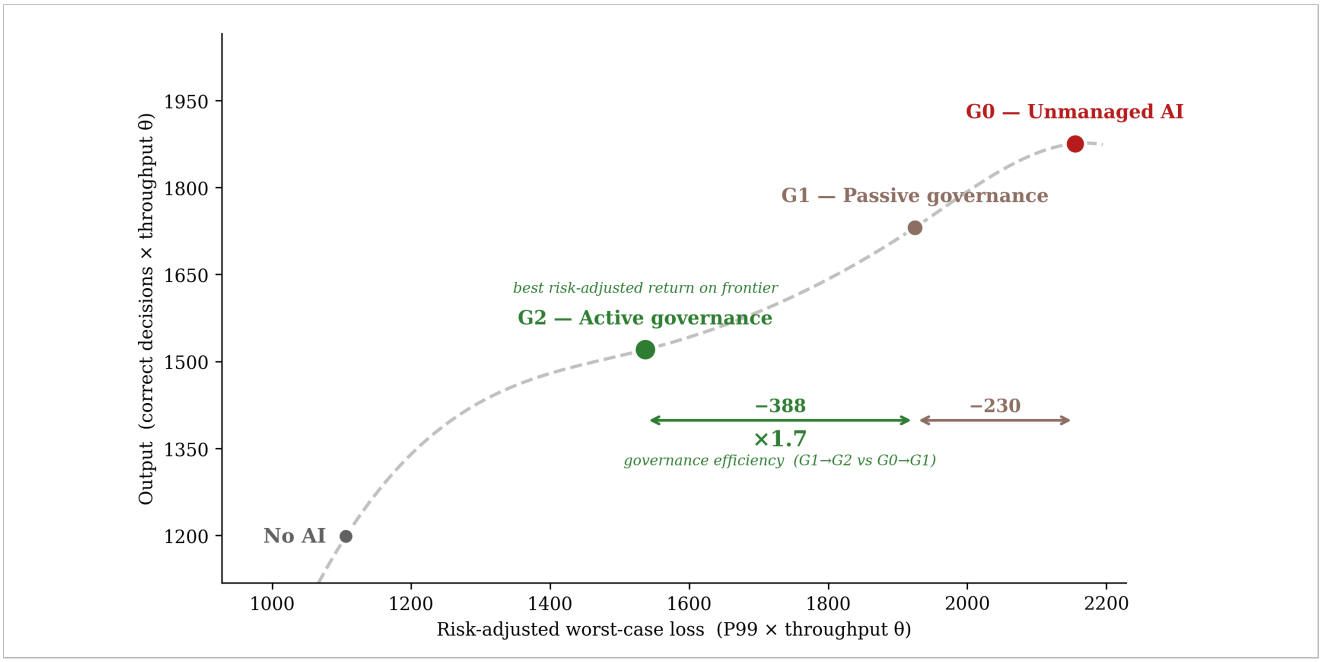

Figure 3. Risk-return frontier across governance regimes. Output versus risk-adjusted $P_{99} \times \theta$. G2 at the elbow: best risk-adjusted return. The first gains (Baseline → G2) are cheap in risk; the next gains (G2 → G0) are expensive.

Active scaffolding protects the cognitive diversity that exists. It cannot restore what biased screening has already eliminated. An annual audit of ATS outputs for cognitive homogenisation is the only intervention that acts on the source. Under the one jurisdiction that requires such audits, 4.6% of organisations comply [35]. Without both levers, the scaffold depletes the resource the hiring audit is meant to preserve.

Neither lever will be voluntarily adopted at scale. The reason is the arithmetic of who absorbs the risk.

Tail risk in the model does not depend on firm size. The fraction absorbed does.

A startup that fails externalises its losses to clients and investors. With a five-year survival probability near 10%, the expected cost of tail risk is negligible against the 54% productivity advantage of unmanaged adoption. The startup rationally chooses G0. Incumbents face pressure to match or concede market share. The Nash equilibrium without regulation is unmanaged adoption for all firms, individually rational,

Firm type	Absorbed risk	Perceived P_{99}
AI-native startup	∼10%	∼171
Mid-market firm	∼60%	∼1,029
Corporation	∼100%	∼1,714

collectively catastrophic. This is a structural argument, not a formally calibrated game; the direction follows from the model's output differentials.

In finance, this diagnosis was made after 2008 and acted upon. The ESRB now identifies model uniformity and provider concentration as systemic vulnerabilities [36]. The FSB warns on third-party dependencies and service-provider concentration [37]. The BIS documents the narrowing of the provider base [38]. The diagnosis is converging. The structural response has not yet arrived outside finance. Under the only jurisdiction requiring AI hiring audits, compliance stands at 4.6% [35]. The EU AI Act enters force in August 2026 without provider concentration limits. Its human-oversight requirements address individual decisions; they do not address the portfolio-level correlation that emerges when thousands of organisations use the same model simultaneously.

The regulatory logic is identical to Basel capital requirements. Banks hold capital not because they are individually irrational but because the unregulated equilibrium is collective overleveraging. AI governance requires the same structure: mandatory provider diversification limits, minimum decision-audit standards, and liability frameworks that internalise tail risk for firms of all sizes. The stack is the natural regulatory target. A regulator can impose concentration limits on AI providers exactly as Basel imposes concentration limits on credit exposures. The scaffold touches internal processes that regulators cannot easily mandate or verify.

Three actors can act without waiting for regulation. One requires it.

Actor	Lever	Action
CRO / board	Monitoring	Error-correlation KRI; concentration-adjusted tail exposure
Procurement lead	Stack architecture	Architectural diversification by layer, actionable in weeks
COO / BU leaders	Process design	Active scaffold: independent judgment before AI consultation
CHRO	Hiring governance	Annual ATS audit for cognitive homogenisation bias
Regulator	Structural framework	Provider concentration limits (Basel model); minimum audit standards

The case for acting now is not only the magnitude of the current exposure. It is that the same channels creating the risk also make governance progressively more expensive. The next section explains why.

6 Why waiting makes governance harder

Three constraints bound the model's scope. The link between cognitive-profile diversity and error decorrelation is an inference. Keck and Tang measure cognitive processes in controlled experiments, not recruitment pipelines in operating organisations [26]. The simulation is a stress test, not a forecast: it demonstrates that the mechanism can produce the paradox under plausible parameters, not that it does so in every organisation. Social influence is not uniformly harmful. Becker, Brackbill and Centola show that decentralised information exchange improves collective accuracy [39]. AI reverses the conditions under which influence is beneficial by centralising the signal and eliminating independence.

These limitations all point in the same direction: the model understates the risk it documents. What it omits (the feedback loop from deskilling to surrender, inter-firm labour market contagion, internal promotion effects that compound homogeneity over time, the scaling of G0-native startups into larger structures) operates through the same mechanism, at the scale of the whole economy.

The most consequential omission is the depletion of a collective resource that no single organisation has incentive to preserve.

Each firm that maximises its AI productivity depletes what we term the cognitive commons: the diversity of approaches, judgments, and error patterns that protects against systemic failure. The structural conditions of Hardin's tragedy [40] are satisfied. There is a non-rival, non-excludable public good, a negative externality, and no individual incentive to preserve the resource. The analogy is not exact, but the arithmetic is suggestive. In monoculture farming, catastrophic crop failure strikes every eight years on average, against once per century in polyculture [41]. The degradation is invisible, masked by the productivity gain that accumulates alongside it. And it is irreversible in a way that Hardin's pasture is not. Deskilling destroys the option to reverse course: endoscopists exposed to AI for diagnostic support showed measurable skill degradation that did not recover when the AI was removed [42]; AI use simultaneously degrades the metacognitive capacity to recognise that degradation [43]. Employees who have not exercised independent judgment for five years cannot resume it on demand. The competence required to revert is the same competence the degradation erodes.

Each quarter of G0 depletes the diversity that makes G2 effective and raises the cost of transition for the next quarter.

The Nash equilibrium hardens because all three channels simultaneously erode the resource and the capacity to restore it. Screening eliminates the cognitively diverse profiles that would make scaffolding effective. Deskilling degrades the independent competence that scaffolding requires. Market competition punishes firms that govern while rewarding those that do not. Beyond individual organisations, a composition effect operates at the economy level. AI-native startups that scale, or are acquired, carry their unmanaged practices into larger structures, gradually displacing the baseline of the whole economy toward G0. The organisations hardest to govern are also the fastest-growing, and their practices do not stay contained.

Whether organisations will need critical human judgment in ten years is not in doubt. Whether the workforce will still be able to exercise it is.

References

[1] L. Meincke, G. Nave, and C. Terwiesch, "ChatGPT decreases idea diversity in brainstorming," *Nature Human Behaviour*, vol. 9, no. 6, pp. 1107–1109, 2025.

[2] S. E. Page, *The Difference: How the Power of Diversity Creates Better Groups, Firms, Schools, and Societies*. Princeton University Press, 2007.

[3] A. Krogh and J. Vedelsby, "Neural network ensembles, cross validation, and active learning," in *Advances in Neural Information Processing Systems*, vol. 7, pp. 231–238, 1995.

[4] D. Wood, T. Mu, A. M. Webb, H. W. J. Reeve, M. Luján, and G. Brown, "A unified theory of diversity in ensemble learning," *Journal of Machine Learning Research*, vol. 24, pp. 1–49, 2023.

[5] N. Holzner, S. Maier, and S. Feuerriegel, "Generative AI and creativity: A systematic literature review and meta-analysis," 2025. Preprint.

[6] J. M. Carlson and J. Doyle, "Highly optimized tolerance: Robustness and design in complex systems," *Physical Review Letters*, vol. 84, no. 11, pp. 2529–2532, 2000.

[7] J. B. Fuller, M. Raman, E. Sage-Gavin, and K. Hines, "Hidden workers: Untapped talent," tech. rep., Harvard Business School and Accenture, 2021.

[8] J. An, D. Huang, C. Lin, and M. Tai, "Measuring gender and racial biases in large language models: Intersectional evidence from automated resume evaluation," *PNAS Nexus*, vol. 4, no. 3, p. pgaf089, 2025.

[9] K. Glazko, Y. Mohammed, B. Kosa, V. Potluri, and J. Mankoff, "Identifying and improving disability bias in GPT-based resume screening," in *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, 2024.

[10] K. Wilson and A. Caliskan, "Gender, race, and intersectional bias in resume screening via language model retrieval," in *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2024.

[11] L. A. Rivera, "Hiring as cultural matching: The case of elite professional service firms," *American Sociological Review*, vol. 77, no. 6, pp. 999–1022, 2012.

[12] B. Schneider, "The people make the place," *Personnel Psychology*, vol. 40, no. 3, pp. 437–453, 1987.

[13] E. M. Kim, A. Garg, K. Peng, and N. Garg, "Correlated errors in large language models," in *Proceedings of the 42nd International Conference on Machine Learning*, vol. 267 of *PMLR*, pp. 30038–30066, 2025. arXiv:2506.07962.

[14] P. Jorion, "Risk management lessons from long-term capital management," *European Financial Management*, vol. 6, no. 3, pp. 277–300, 2000.

[15] D. MacKenzie, "Long-term capital management and the sociology of arbitrage," *Economy and Society*, vol. 32, no. 3, pp. 349–380, 2003.

[16] A. G. Haldane and R. M. May, "Systemic risk in banking ecosystems," *Nature*, vol. 469, no. 7330, pp. 351–355, 2011.

[17] Microsoft Security Response Center, "Helping our customers through the CrowdStrike outage." Microsoft Blog, July 2024. `https://blogs.microsoft.com/blog/2024/07/20/helping-our-customers-through-the-crowdstrike-outage/`.

[18] Parametrix Solutions, "CrowdStrike's impact on the Fortune 500," tech. rep., Parametrix, 2024.

[19] D. Geer, R. Bace, P. Gutmann, P. Metzger, C. P. Pfleeger, J. S. Quarterman, and B. Schneier, "CyberInsecurity: The cost of monopoly," tech. rep., Computer & Communications Industry Association, 2003.

[20] J. M. Kleinberg and M. Raghavan, "Algorithmic monoculture and social welfare," *Proceedings of the National Academy of Sciences*, vol. 118, no. 22, p. e2018340118, 2021.

[21] S. D. Shaw and G. Nave, "Thinking—fast, slow, and artificial: How AI is reshaping human reasoning and the rise of cognitive surrender," 2026. PsyArXiv preprint.

[22] F. Dell'Acqua, E. McFowland III, E. R. Mollick, H. Lifshitz-Assaf, K. Kellogg, S. Rajendran, L. Krayer, F. Candelon, and K. R. Lakhani, "Navigating the jagged technological frontier: Field experimental evidence of the effects of AI on knowledge worker productivity and quality," *Organization Science*, 2026. Articles in Advance.

[23] S. Jabbour, D. Fouhey, S. Shepard, T. S. Valley, E. A. Kazerooni, N. Banovic, J. Wiens, and M. W. Sjoding, "Measuring the impact of AI in the diagnosis of hospitalized patients: A randomized clinical vignette survey study," *JAMA*, vol. 330, no. 23, pp. 2275–2284, 2023.

[24] A. D. Bryant, G. S. Fletcher, and T. H. Payne, "Drug interaction alert override rates in the meaningful use era: No evidence of progress," *Applied Clinical Informatics*, vol. 5, no. 3, pp. 802–813, 2014.

[25] L. Bainbridge, "Ironies of automation," *Automatica*, vol. 19, no. 6, pp. 775–779, 1983.

[26] S. Keck and W. Tang, "Enhancing the wisdom of the crowd with cognitive-process diversity: The benefits of aggregating intuitive and analytical judgments," *Psychological Science*, vol. 31, no. 10, pp. 1272–1282, 2020.

[27] O. A. Vašíček, "Loan portfolio value," *Risk*, vol. 15, no. 12, pp. 160–162, 2002. Original theory in KMV Corporation working paper, 1987.

[28] I. Aggarwal, R. P. DeSimone, F. Kunze, S. Braun, and T. Rigotti, "The dark and the bright side of cognitive diversity: The role of team climate in the relationship between team cognitive diversity and team performance," *Frontiers in Psychology*, vol. 10, p. 112, 2019.

[29] S. E. Asch, "Effects of group pressure upon the modification and distortion of judgments," in *Groups, Leadership, and Men: Research in Human Relations* (H. Guetzkow, ed.), pp. 177–190, Pittsburgh, PA: Carnegie Press, 1951.

[30] Menlo Ventures, "2025: The state of generative AI in the enterprise," tech. rep., Menlo Ventures, Dec. 2025. `https://menlovc.com/perspective/2025-the-state-of-generative-ai-in-the-enterprise/`.

[31] S. Randazzo, A. Joshi, K. Kellogg, H. Lifshitz-Assaf, F. Dell'Acqua, and K. R. Lakhani, "GenAI as a power persuader: How professionals get persuasion bombed when they attempt to validate LLMs," Working Paper 26-021, Harvard Business School, 2025.

[32] M. von Zahn, L. Liebich, E. Jussupow, O. Hinz, and K. Bauer, "Knowing (not) to know: Explainable artificial intelligence and human metacognition," *Information Systems Research*, 2025.

[33] V. Frey and A. van de Rijt, "Social influence undermines the wisdom of crowds in sequential decision making," *Management Science*, vol. 67, no. 7, pp. 4303–4316, 2021.

[34] A. B. Haynes, T. G. Weiser, W. R. Berry, S. R. Lipsitz, A.-H. S. Breizat, E. P. Dellinger, *et al.*, "A surgical safety checklist to reduce morbidity and mortality in a global population," *New England Journal of Medicine*, vol. 360, no. 5, pp. 491–499, 2009.

[35] L. Wright, R. M. Muenster, B. Vecchione, T. Qu, S. Cai, J. Metcalf, and J. N. Matias, "Null compliance: NYC local law 144 and the challenges of algorithm accountability," in *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, 2024.

[36] European Systemic Risk Board, Advisory Scientific Committee, "Artificial intelligence and systemic risk in the financial system," ASC Report 16, ESRB, 2025. `https://www.esrb.europa.eu/pub/pdf/asc/esrb.ascreport202512_AIandsystemicrisk.en.pdf`.

[37] Financial Stability Board, "Monitoring the adoption of artificial intelligence and related vulnerabilities in the financial sector," tech. rep., FSB, Oct. 2025. `https://www.fsb.org/uploads/P101025.pdf`.

[38] Bank for International Settlements, Financial Stability Institute, "Financial stability implications of artificial intelligence," executive summary, BIS/FSI, June 2025. `https://www.bis.org/fsi/fsisummaries/exsum_23904.htm`.

[39] J. Becker, D. Brackbill, and D. Centola, "Network dynamics of social influence in the wisdom of crowds," *Proceedings of the National Academy of Sciences*, vol. 114, no. 26, pp. E5070–E5076, 2017.

[40] G. Hardin, "The tragedy of the commons," *Science*, vol. 162, no. 3859, pp. 1243–1248, 1968.

[41] D. Renard and D. Tilman, "National food production stabilized by crop diversity," *Nature*, vol. 571, no. 7764, pp. 257–260, 2019.

[42] K. Budzyń, M. Romańczyk, D. Kitala, P. Kołodziej, M. Bugajski, H. O. Adami, J. Blom, M. Buszkiewicz, N. Halvorsen, C. Hassan, T. Romańczyk, Ø. Holme, K. Jarus, S. Fielding, M. Kunar, M. Pellise, N. Pilonis, M. F. Kamiński, M. Kalager, M. Bretthauer, and Y. Mori, "Endoscopist deskilling risk after exposure to artificial intelligence in colonoscopy: A multicentre, observational study," *The Lancet Gastroenterology & Hepatology*, vol. 10, no. 10, pp. 896–903, 2025.

[43] Y. Fan, L. Tang, H. Le, K. Shen, S. Tan, Y. Zhao, Y. Shen, X. Li, and D. Gašević, "Beware of metacognitive laziness: Effects of generative artificial intelligence on learning motivation, processes, and performance," *British Journal of Educational Technology*, 2025.